

MLPM project

The effect of smoothness

Maarten van der Velden
5743087
`Maarten.vanderVelden@student.uva.nl`

Carsten van Weelden
0518824
`cweelden@science.uva.nl`

November 4, 2011

1 Introduction

In this paper we investigate the effect of smoothness of the dataset on the performance of classification algorithms. In order to investigate this we generate several artificial datasets of varying smoothness and look at the accuracy and loss of the resulting classifiers. In order to get insight into the effects we decompose the loss into its bias, variance, and noise components as defined in [?].

To investigate the effect that smoothness has on classification performance we first need to define some idea of smoothness which we can easily vary and then run a set of experiments for different levels of smoothness. We view classification as a two class problem¹, with each class being represented by some Probability Density Function (PDF) over the attribute space. Given this view the *smoothness* of a class is determined by the shape of the PDF. We define the PDF for each class as a Gaussian Mixture Model (GMM) with k mixture components. Intuitively, the smoothness is then determined by the number of components in the mixture. With just one component the PDF corresponds to a Gaussian distribution which is very smooth, while increasing the number of components increases the peakedness of the distribution, making it less smooth.

One way to see this is in terms of inherent noise in the problem, which we define as the noise in definition 4 of [?]:

$$N(x) = E_t[L(t, y_*)] \quad (1)$$

If we keep the mean variance of the distributions equal, then with more components we will have a higher average probability that the PDFs overlap for a datapoint x . Since the Bayes optimal prediction predicts the class for which the PDF is highest at x , the noise is proportional to the area under the PDF

¹As contrasted with a concept learning problem, in which there is one class which needs to be distinguished from the background.

for which the PDF of the other class is higher. In other words, the more components make up the distribution of a class, the more overlap there is between classes, and therefore the higher the noise component in the loss is.

To show this effect we generated GMMs with an increasing amount of components and measured their noise. Averaged over 25 random models of the same number of components, it is clear that the noise increases with the number of components, growing asymptotically towards 0.5, which represents the amount of noise where the Bayes optimal decision is correct half of the time: chance level. The results are shown in Figure 1.

2 Experimental method

1. We generate a problem containing two classes denoted C_0 and C_1 . Each class is represented by a PDF which we define as a GMM with k components². We do this for $k = [1..5]$. We pick the parameters corresponding to component j of class i as follows. The prior probability for each component is $\pi_{ij} = \text{rand}([1..5])/Z_i$ where Z_i is a normalizing factor such that $\sum_j \pi_{ij} = 1$. The mean of each component is $\mu_{ij} = (\text{rand}([-1, 1]), \text{rand}([-1, 1]))$. We use a scalar covariance matrix σI with $\sigma = \text{rand}([0.1, 0.4])$.
2. For each two class problem we generate a set of 100 training sets $D = \{d_1, d_2, \dots, d_{100}\}$ and a corresponding test set T . We do this for $|D| \in 10, 100, 1000, 1000$ with half of the data points being sampled from each class. For the test set $|T| = 1000$, also with half of the points being sampled from each class.
3. We train a classifier on each training set $d \in D$ and evaluate the accuracy on the test set.
4. Finally, we compute the average bias and average variance on T .

3 Results

3.1 Naive Bayes

We trained Naive Bayes classifiers for all settings of the GMM we defined in the last section. The resulting error is shown in Figure 2, and the subdivision of this into bias and variance can be seen in Figures 3 and 4 respectively.

The figures show that in general, less training data results in a higher loss and variance, which was to be expected. Furthermore, the loss is higher when the smoothness increases. This might result solely from there being more noise, but Figure 3 shows that the bias increases sharply on less smooth data too. There is less increase in the amount of variance overall, but this increase seems to be more significant with less training data than on much training data.

3.2 kNN

4 Conclusion

²We use a 2-dimensional attribute space for ease of visualization.




Figure 1: The effect on the noise of the use of more mixture components in the GMM.




Figure 2: The average error of the Naive Bayes classifier on different amounts of training data, given different levels of smoothness of the data.




Figure 3: The average bias of the Naive Bayes classifier on different amounts of training data, given different levels of smoothness of the data.




Figure 4: The average variance of the Naive Bayes classifier on different amounts of training data, given different levels of smoothness of the data.